

# COMPSCI361: Machine Learning

## Data Preprocessing

Katerina Taskova and Jörg Simon Wicker  
The University of Auckland



**SCIENCE**  
SCHOOL OF COMPUTER SCIENCE

## Data Preprocessing

# This lecture will cover



SCIENCE  
SCHOOL OF COMPUTER SCIENCE

## Data Preprocessing

### Data Reduction

- Dimensionality Reduction

- Principal Components Analysis

- Feature Selection

## Data Reduction

# Data Reduction

- Data reduction: Obtain a reduced representation of the data set that is much smaller in volume but yet produces the same (or almost the same) analytical results
- Why data reduction? – A database may store terabytes of data, complex data analysis may take a very long time to run on the complete data set
- Data reduction strategies
  - Dimensionality reduction
    - Wavelet transforms
    - Principal Components Analysis (PCA)
    - Feature selection
  - Numerosity reduction
    - Regression and Log-Linear Models
    - Histograms, clustering, sampling
  - Data compression

# Dimensionality Reduction



## ■ Curse of dimensionality

| $x_1$ | $x_2$ | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | $x_n$ |
|-------|-------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-------|
| 0     | 0     | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 7   | 11  | 12  | 22  | 24    |
| 0     | 0     | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 74  | 38  | 99  | 2   | 4   | 0   | 0   | 0   | 0   | 0     |
| 0     | 0     | 0   | 0   | 0   | 0   | 0   | 0   | 84  | 69  | 55  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0     |
| 0     | 0     | 0   | 0   | 66  | 35  | 14  | 62  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0     |
| 32    | 48    | 54  | 21  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0     |
| ...   | ...   | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ...   |


- When dimensionality increases, data becomes increasingly sparse
- Density and distance between points, which is critical to clustering, outlier analysis, classification, regression becomes less meaningful

# Dimensionality Reduction

- Why dimensionality reduction?
  - Avoid the curse of dimensionality
  - Help eliminate irrelevant/redundant features and reduce noise
  - Reduce computational resources (memory and time)
  - Allow easier visualization
- Dimensionality reduction techniques
  - Unsupervised linear method: Principal Component Analysis
  - Supervised method: Feature selection

# Principal Component Analysis – PCA

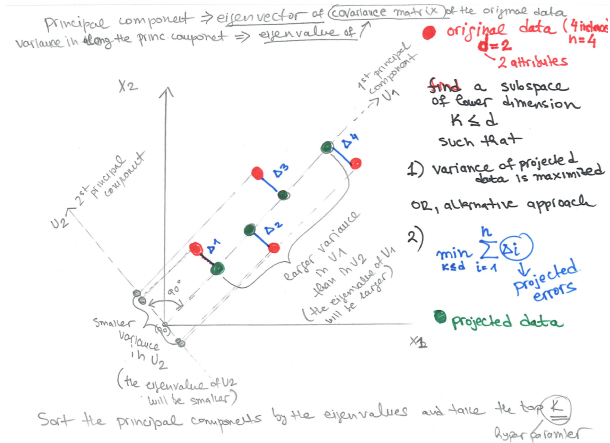
- Find a projection that captures the largest amount of variation in data
- The original data are projected onto a much smaller space
- How? – We find the eigenvectors and eigenvalues of the covariance matrix of the input attributes,
  - **Eigenvectors:** the directions that data variances occur, and they define the new attribute space
  - **Eigenvalues:** the amount of variance along the corresponding eigenvector



Demo: <https://setosa.io/ev/principal-component-analysis/>



# PCA approaches



## PCA – steps

- Given  $n$  data instances (each a vector in  $d$ -dimensions), find  $k \leq d$  principal components that can be best used to represent the data
  - Normalize input data: Each attribute falls within the same range
  - Compute  $k$  orthonormal vectors (i.e. length of 1 and perpendicular to each other), i.e. principal components
  - The input data is a linear combination of the  $k$  principal component vectors
  - The principal components are sorted in order of decreasing “significance” or strength (e.g. as measured by the eigenvalue)
  - Since the components are sorted, the size of the data can be reduced by eliminating the  $d - k$  weak components, i.e. those with low variance.
- The resulting vectors are orthogonal, are they correlated? No
- Can you use PCA on categorical data? Yes

Example code step by step PCA calculation in Python in reference 3

# Feature or Attribute Selection

Reduce dimensionality by removing set of attributes

- **Redundant attributes**

- Duplicate much or all of the information contained in one or more other attributes  
e.g. purchase price of a product and the amount of sales tax paid

- **Irrelevant attributes**

- Contain no information that is useful for the data mining task at hand  
e.g. students' ID is often irrelevant to the task of predicting students' GPA

Two types of methods: **Filters** (fast) and **Wrappers** (high accuracy, expensive)

- Filters separate feature selection from classifier learning. No bias toward any learning algorithm

## Feature Selection using Correlation

- For nominal data, given two attributes  $A$  and  $B$  with values  $a_1, \dots, a_c$  and  $b_1, \dots, b_r$  the correlation can be calculated using the  $\chi^2$  test:

$$\chi^2 = \sum_{i=1}^c \sum_{j=1}^r \frac{(o_{ij} - e_{ij})^2}{e_{ij}}$$

- With  $o_{ij}$  being the actual frequency of the event  $(a_i, b_j)$
- And  $e_{ij}$  the expected frequency ( $n$  is the number of instances)

$$e_{ij} = \frac{\text{count}(A = a_i) \times \text{count}(B = b_j)}{n}$$

## Feature Selection using Correlation


- Numerical data can be compared using Pearson's correlation coefficient

$$\rho_{A,B} = \frac{\sum_{i=1}^n (a_i - \bar{A})(b_i - \bar{B})}{n\sigma_A\sigma_B} = \frac{\sum_{i=1}^n (a_i b_i) - n\bar{A}\bar{B}}{n\sigma_A\sigma_B}$$

- With means  $\bar{A}$  and  $\bar{B}$ , number of instances  $n$ , and standard deviations  $\sigma_A$  and  $\sigma_B$
- If  $\sigma_A$  and  $\sigma_B$  is zero then the coefficient is undefined.
- Values in the range  $[-1, 1]$ .

# Feature Selection using Correlation

- So what does the correlation measure?



By Kiatdd - Own work, CC BY-SA 3.0,

[commons.wikimedia.org/w/index.php?curid=37108966](https://commons.wikimedia.org/w/index.php?curid=37108966)

- How can it be used to remove redundant or unimportant features?

# Heuristic Search in Attribute Selection

- There are  $2^{d-1}$  possible attribute combinations of  $d$  attributes  
→ exhaustive search is not feasible (e.g.  $d = 300$ ,  $2.04 \times 10^{90}$  combinations)
- Typical heuristic attribute selection methods:
  - Best single attribute under the attribute independence assumption
  - Best step-wise feature selection:
    - The best single-attribute is picked first
    - Then next best attribute condition to the first, ...
  - Step-wise attribute elimination:
    - Repeatedly eliminate the worst attribute
  - Best combined attribute selection and elimination
  - Optimal branch and bound:
    - Use attribute elimination and backtracking

# Relief - Instance-based heuristic for feature selection

**Input:** Data set with  $n_d$  input attributes and  $n$  instances that belong to one of two classes (i.e. binary classification problem), and number of randomly selected instances  $n_r \leq n$

First normalize the input attributes

Create a weight vector  $W$  with one weight  $w_i \in W$  for each attribute

Initialize the weights  $W = [w_1, w_2, \dots, w_n] = 0$

**for**  $j \in 1 \dots n_r$  **do**

    Randomly select instance  $R = [r_1, \dots, r_n]$

    Choose instance  $H = [h_1, \dots, h_n]$  as the closest neighbour of  $R$  in the same class (*nearHit*) w.r.t to some (Euclidian) distance measure

    Choose instance  $M = [m_1, \dots, m_n]$  as the closest neighbour of  $R$  in the other class (*nearMiss*) w.r.t to some (Euclidian) distance measure

**for**  $i \in 1 \dots n_d$  **do**

            % update the weights  $w_i = w_i + \text{distance}(R, M; \text{i-th attribute}) - \text{distance}(R, H; \text{i-th attribute})$

$w_i = w_i + (r_i - m_i)^2 - (r_i - h_i)^2$  % Euclidian distance

**end**

**end**

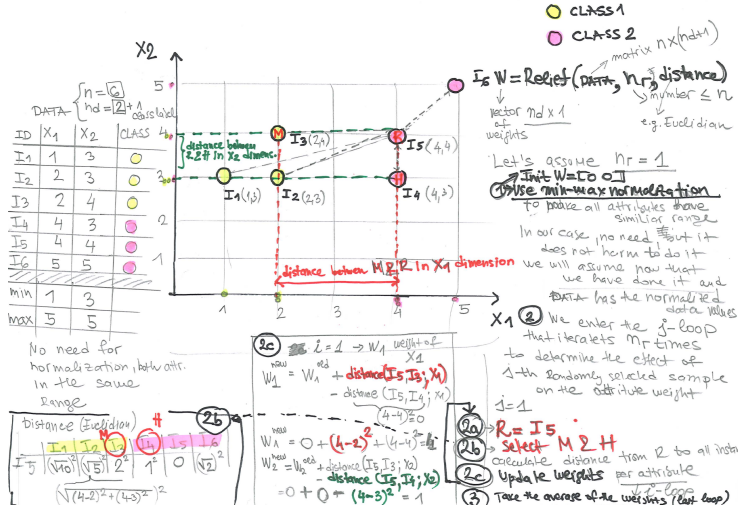
**for**  $i \in 1 \dots n_d$  **do**

    return  $w_i = \frac{w_i}{n_r}$

**end**



# Relief Example



## Relief Example (cont.)

In our case  $h_r = 1$ , so no more iterations  
in the  $j$ -loop



We go to the last step ③

$W = \begin{bmatrix} 4 & 1 \end{bmatrix}$  before step ③

$W = \begin{bmatrix} 4 & 1 \\ 1 & 1 \end{bmatrix}$  after step ③  $\rightarrow$  no change as only  
one random sample was selected

You can use these weights to rank the attributes  
Larger weight more important

In our case  $X_1$  is the more important

Easy to verify: from the plot we  
can see that we can use only  
 $X_1$  to create decision boundary  
that will separate the  samples  
from the  samples

Also note that the variance in the data  
is larger in the  $X_1$  dimension (link to PCA  
2 projections appearance)


## Relief summary



- Relief takes into account **all** attributes
- Result is a weight vector that represents the importance of each feature
- Features are then selected based on a threshold  $\tau$  or ranked
- The algorithm above is the basic version of Relief, there are various extensions (ReliefF, RReliefF, ...)

# Wrappers

- The correlation method and Relief are **filters**
- Main idea of **wrappers**:
  - Generate a subset of the features and evaluate the performance of the classifier on the subset
  - Add or remove attributes from the subset and see if the performance of the classifier improves
  - Risk of overfitting, especially if choosing the same classifier as for the main learning task



# Preprocessing



- So, to summarize...
  - When are preprocessing approaches useful?
  - When should you avoid them?
  - How about specific cases
    - Many correlated features?
    - Many independent features?
    - Which algorithms you know already would need preprocessing?
    - How about Decision trees? Why?
    - How about Regression? Why?
  - Are we cheating in preprocessing: for example by creating new examples?

## Conclusion

- Preprocessing is an important part in machine learning and data analysis
- Missing values can be caused by various reasons depending on what the reasons are, they must be addressed differently
- Various imputation approaches exist, they use the information of other instances and values to impute the missing values
- Noisy data can be addressed for example by binning, clustering, or regression
- Feature selection can be used to reduce the number of redundant and unimportant features
- Imbalanced data sets can be a problem for evaluation and classifiers
- Sampling can be used to overcome class imbalance problems

## Literature & other resources

1. Material in Chapter 3 in Han's *Data Mining*
2. Detailed math and application of PCA: Chapter 12.1 and Appendix C Bishop's *Pattern Recognition and Machine Learning*
3. Step by step code on PCA calculation (Python notebook on Canvas)
4. Example code on using scikit-learn PCA implementation  
<https://colab.research.google.com/github/cpearce/PythonDataScienceHandbook/blob/master/notebooks/05.09-Principal-Component-Analysis.ipynb>

Thank you for your attention!

`https://ml.acukland.ac.nz`