



COMPSC 361

Advanced Neural Networks

Instructor : Thomas Lacombe

Neural Networks III
Week 11

Disclaimer/Acknowledgment:

The following slides are reusing some of the content created by Andrew Ng in his book “Machine Learning Yearning” and his Coursera course about Improving DNN.

<https://github.com/daiwk/ml-yearning/blob/master/Ng-MLY01-13.pdf>

<https://www.deeplearning.ai/courses/deep-learning-specialization/>



Outline

Introduction

Artificial Neural Networks (ANN)

- Single Unit: Architecture of Perceptron (NN1)
- Connection to Shallow Machine Learning (NN1)
- Multi-Layer Feed-Forward Neural Network (NN2)

Design Issues (NN3)

Deep Learning / Large Language Models (NN4)



Empirical and iterative process

Neural network design:
Lots of choices to make!

- ▶ Evaluation:

Cost function

Evaluation strategy

- ▶ NN hyperparameters:

Number of layers

Number of neurons per layer


Activation functions

Learning rate

Weight initialisation

Mini-batch size

...





Design issues: Evaluation strategy

▶ Train/dev/test sets with deep networks



- ▶ If small dataset (100, 1000, 10 000 samples)

↳ 60%/20%/20%


- ▶ If large dataset (> 1 000 000)

↳ 98%/1%/1%


- ▶ Training set and dev/test set usually need to come from same distribution (but it is ok if it varies a bit when gathering a lot of training data).
- ▶ Make sure dev and test sets come from the same distribution.

How does your model do?


► Bias vs variance / underfitting vs overfitting






high bias
Underfit




“just right”



high variance
Overfit




How does your model do?




training error = 15%
dev error = 16%

High bias
Underfitting




training error = 3%
dev error = 4%

Low bias
Low variance
Good performance



training error = 3%
dev error = 15%

High variance
Overfitting




training error = 15%
dev error = 35%

High bias
High variance



How to improve learning? Overfitting

- ▶ Very common problem with Deep Learning: overfitting



- ▶ Regularisation = discouraging learning a more complex model
- ↳ Reduces the variance, but increases the bias



Design issues: Regularisation

“Regularization is any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error.”

Deep learning, 5.2.2, p.117

► Different technics:





- ↪ L1 & L2 regularisation
- ↪ Dropout
- ↪ Early stopping
- ↪ Data augmentation

<https://www.analyticsvidhya.com/blog/2018/04/fundamentals-deep-learning-regularization-techniques/>

How does regularisation help to avoid overfitting?

First intuition:

- ▶ E.g., L1 & L2 regularisation penalises large weights values.
- ▶ Keeping weights close to zero for some neurons.






How does regularisation help to avoid overfitting?

Second intuition:

- ▶ Limiting the weights values will bring the output of a neuron in the linear zone of activation function (e.g. \tanh).
- ▶ It will limit the NN power to model non-linearities.





L1 and L2 regularisation

► Why use it?

Large weights:

- are characteristic of more complex models (higher learning time).
- can be a sign of an over-specialized network (overfitting).
- make the network unstable (sensitive to noise).

Penalises/constrains the weight values towards 0.

↪ A "weight shrinkage" or a "penalty against complexity"

↪ Encourages simpler models.



L1 and L2 regularization

L1 and L2 norms

$$\|\mathbf{w}\|_1 = |w_1| + |w_2| + \dots + |w_N|$$

$$\|\mathbf{w}\|_2 = (|w_1|^2 + |w_2|^2 + \dots + |w_N|^2)^{\frac{1}{2}}$$

► How does it work?

1. Calculate the weights size

- Sum of the absolute values of the weights \rightarrow L1.
- Sum of the squared values of the weights \rightarrow L2.

$$\sum_{i=1}^N |w_i|$$
$$\sum_{i=1}^N w_i^2$$

2. Apply regularisation to the weight update

- L1 regularisation : $Loss = Error(y, \hat{y}) + \lambda \sum_{i=1}^N |w_i|$
- L2 regularisation: $Loss = Error(y, \hat{y}) + \lambda \sum_{i=1}^N w_i^2$

λ controls the penalty

$$0 < \lambda < 1$$

L1 vs L2 regularisation

► Weight update: $w \leftarrow w - \frac{dL}{dw}$

► L1: $Loss = Error(y, \hat{y}) + \lambda \sum_{i=1}^N |w_i|$

► L2: $Loss = Error(y, \hat{y}) + \lambda \sum_{i=1}^N w_i^2$

→ L2 penalises more large weights and less small weights than L1.

→ L1 shrinks weights to 0 while L2 shrinks weights evenly.




Image source: <https://towardsdatascience.com/visualizing-regularization-and-the-l1-and-l2-norms-d962aa769932>

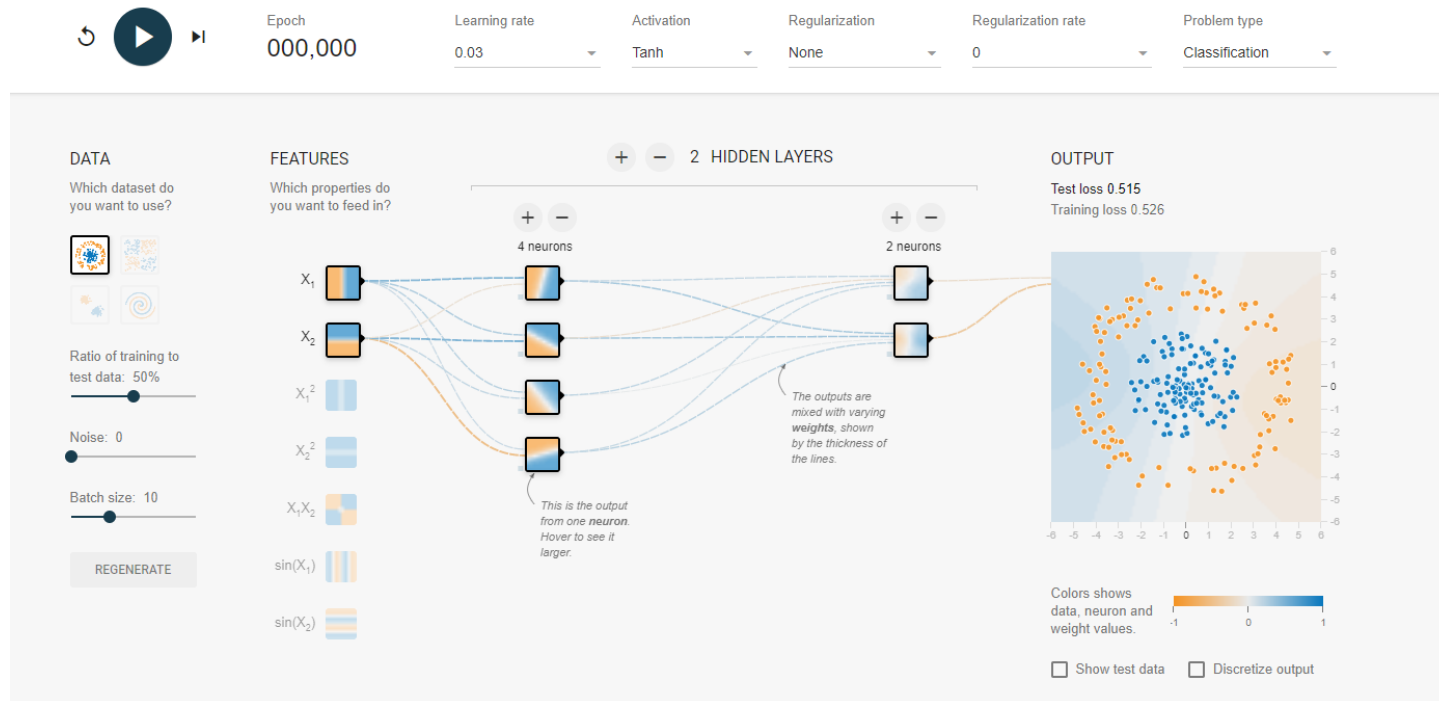


L1 vs L2 regularisation

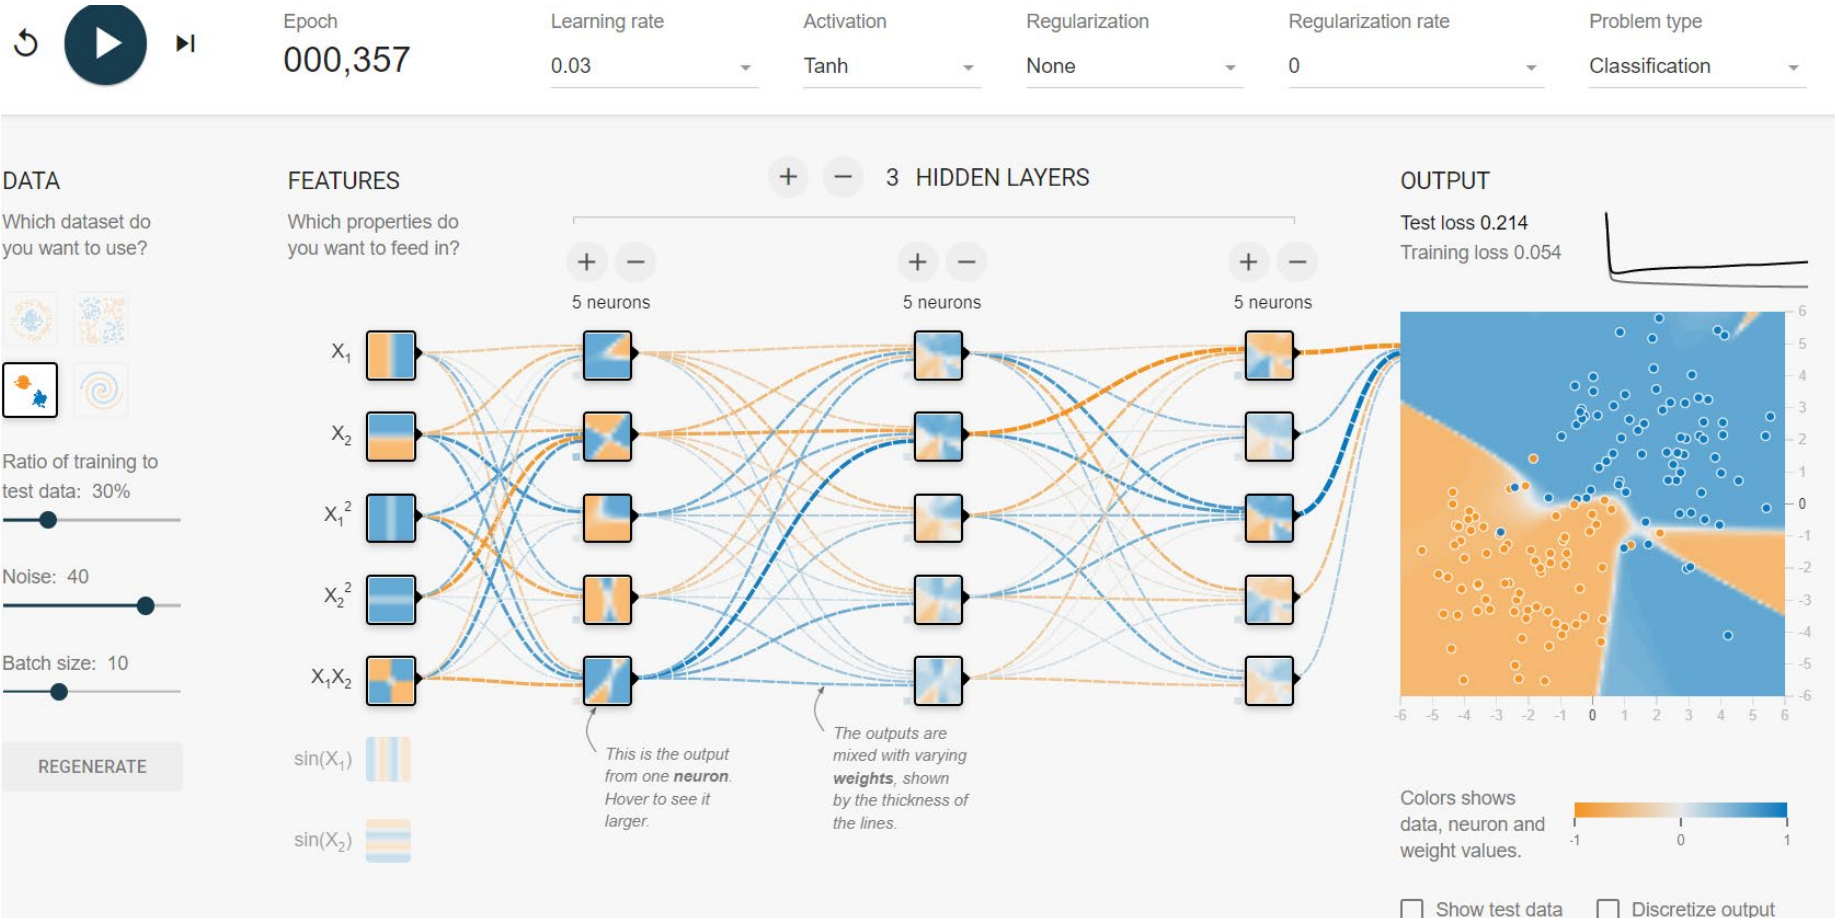
- ▶ L1 regularisation results in a sparse weight matrix (a lot of weight values to 0).
- ▶ L1 regularisation is acting as feature selection, dropping irrelevant features.
- ▶ L2 regularisation results in less sparse weight matrix than L1, and it will reduce the effect of collinear features.
- ▶ Penalising the weights forces the NN to “focus” more on simpler features that explain most of the variance, than on complex ones.

Tensorflow playground platform

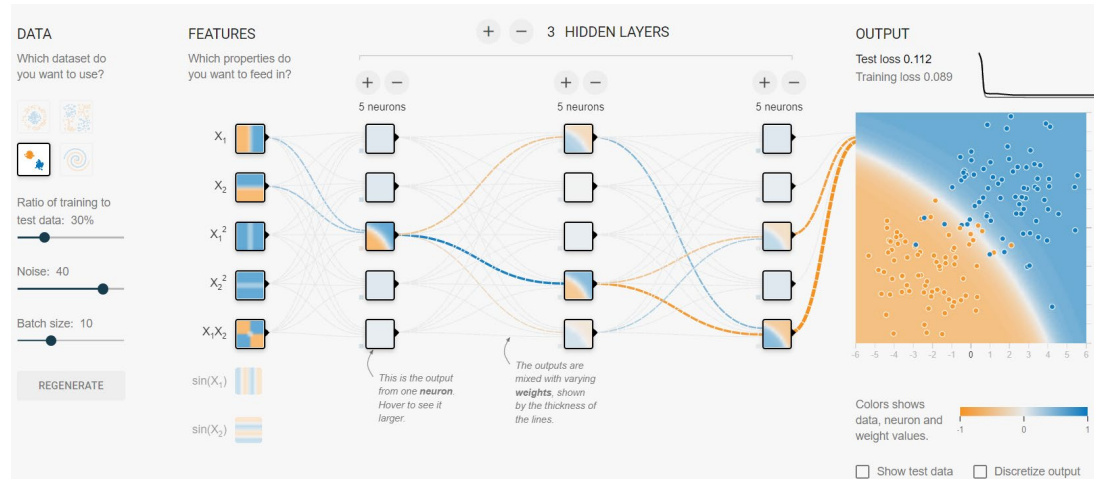
- ▶ Platform to test and visualise the effects of varying hyperparameters: <https://playground.tensorflow.org/>



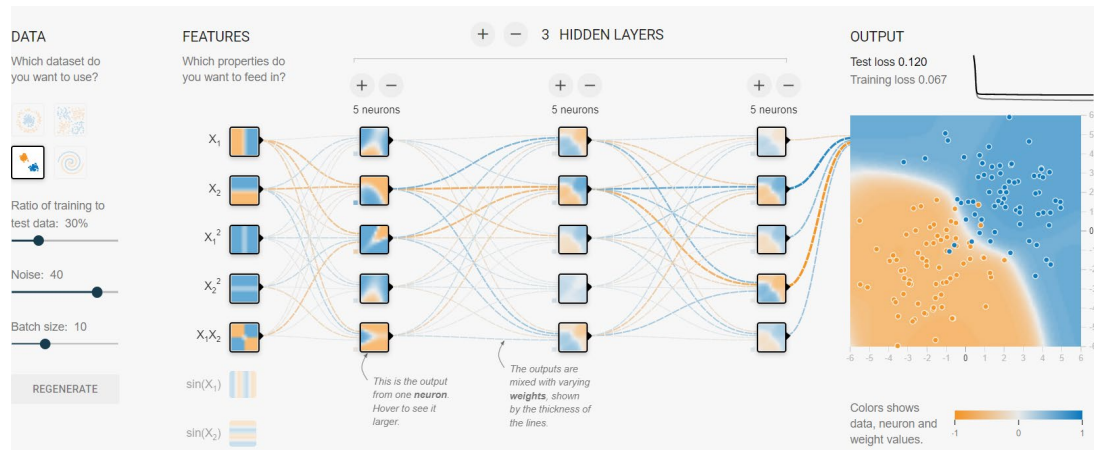
Without regularisation



With regularisation




L1




L2

Dropout

► How does it work?



Probability for hidden layers: $p = 0.5$
Probability for input layer: $0.5 < p < 1.0$



<https://machinelearningmastery.com/dropout-for-regularizing-deep-neural-networks/>
<https://wandb.ai/authors/ayusht/reports/Dropout-in-PyTorch-An-Example--VmldzoxNTgwOTE>
Srivastava, N. et al. (2014). Dropout: a simple way to prevent neural networks from overfitting.



Dropout

► Why does it work?

The nodes cannot rely on any single previous node (feature).


- Prevents too large weights.
- Encourages spreading out the weights.
- Forces the nodes to be more generally useful.

Can be combined with other forms of regularisation.



Early stopping

- ▶ Simple and popular regularisation technic.



- ▶ Learn enough, but not too much!
- ✚ Avoid to end up in the high variance zone.



Early stopping

► When to stop?

1. Monitor the performance

- Loss on the dev dataset.
- Additional metrics (e.g., precision, recall, etc).

2. Trigger the early stopping

- Simplest trigger: increase of the loss compared to the last iterations.
- More elaborated ones: no change over several epochs, absolute change in a metric, average change in a metric over several epochs, reaching a specific level of performance, etc.

3. Choose the model to keep

- Usually, keep the model from the epoch before the increase in loss.

[Prechelt L. \(1998\) Early Stopping - But When?](#)



Regularisation

- ▶ Data augmentation
 - ▶ Overfitting can happen if you do not have enough data to train all parameters.
 - ▶ Pre-processing technic → does not modify explicitly the learning algorithm.
 - ↳ Increases the training data set size.
 - ↳ Increases the diversity of the data.
 - ↳ Especially used with images.
 - ↳ Includes operations like rotating the image, flipping, scaling, adding noise, etc.
 - ⚠ Can lead to underfitting if generated data not relevant to the task.

[Shorten, C., & Khoshgoftaar, T. M. \(2019\). A survey on image data augmentation for deep learning.](#)



Optimal error rate / avoidable bias

► Optimal error rate


- ✚ Error rate of an optimal classifier (e.g., human performance)
- ✚ Can be hard to estimate

► Avoidable bias

- ✚ Training error – optimal error rate

► Variance

- ✚ Dev error – training error





Simplest formula to address variance/bias issues

► High avoidable bias

- ↪ Intuition: model not complex enough to map inputs and outputs.
- ↪ Simple fix: Increase model size (e.g., increase layers or neurons per layer).
- ↪ Might increase variance and risk of overfitting (if no regularisation).
- ↪ Will slow the learning.

► High variance

- ↪ Intuition: training data not sufficient to generalise on dev data.
- ↪ Simple fix: Add data to the training set.
- ↪ More data might not be available.
- ↪ Try data augmentation.



Bias vs variance tradeoff

- ▶ Some choices reduce bias but increase variance.
 - ↪ E.g., increasing size of the network.
- ▶ Some choices reduce variance but increase bias.
 - ↪ E.g., adding regularisation (early stopping might stop learning before reaching low bias, penalizing high weights might prevent the model to reach low bias, etc).
 - ↪ Effect of regularisation on bias can be reduced with a good hyperparameter tuning.
 - ↪ Data augmentation does not increase bias if relevant augmentation.
- ▶ More useful advice in Sections 25 to 27 of Andrew Ng's "Machine Learning Yearning" book, to reduce variance and bias.



Design issues: Initialisation

Initialise the weights and biases in the network

- ▶ Random: E.g., weights are initialized randomly from $\text{Uniform}[-0.1, 0.1]$. Biases are initialized to 0s.
- ▶ Zero: All weights are initialized to 0.
- ▶ With deep networks, always initialise weight randomly (e.g. standard normal distribution) to break the « symmetry ».
- ▶ Additional tip: Also good to normalize inputs to mean zero and use random weight initialization with avg. weight centered at zero.



Design issues: learning rate

Gradient descent is an optimization algorithm that finds the local minimum of a function by taking "**steps**" in the direction of the negative of the gradient.

- ▶ What will happen if we use a learning rate that is too small or too large?
- ▶ Learning efficiency, optimization accuracy
- ▶ E.g., learning steps that were taken to find the local minimum of a function

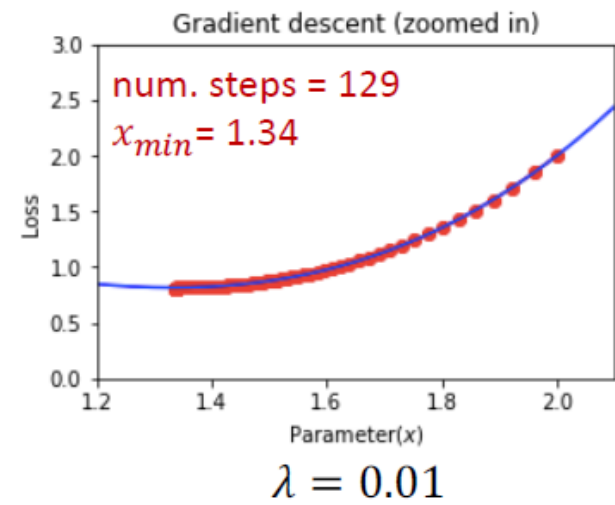




Image source: Meng-Fen Chiang



Design issues: learning rate

► Learning rate decay schedule

- ↪ Common practice: decrease learning rate over time (learning rate decay).
- ↪ E.g., linear decay.
- ↪ Linear decay for set number of iterations and then constant.

► Adaptive learning rates strategies

- ↪ Monitors the model's performance and adapt the learning rate in response.
- ↪ Reduces learning rate when performance plateaus.
- ↪ Increases learning rate when performance does not improve for a number of iterations.



Design issues :Exploding and vanishing gradients

- ▶ Gradients are calculated in the backpropagation process to update the weights in the desired direction.
- ▶ **Vanishing gradients:**
 - ▶ Gradients become smaller and smaller and can become close to 0.
 - ▶ Can slow down or stop the learning process (very small weights' update).

$$\text{out_}z_j = g(\sum_i w_{i,j} x_i + b_j)$$

$$\text{out_}z'_j = g'(\sum_i w_{i,j} x_i + b_j) * x_i \text{ (chain rule)}$$

If $g'()$ is close to 0, then the value of the gradient becomes smaller and smaller as backpropagation processes back to the initial layers (significant for large NN).



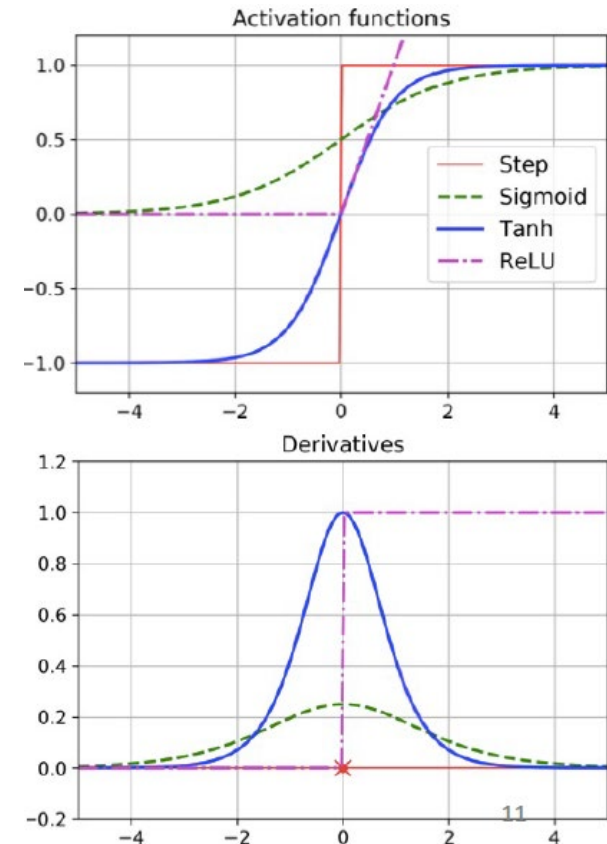
Design issues :Exploding and vanishing gradients

- ▶ Gradients are calculated in the backpropagation process to update the weights in the desired direction.
- ▶ **Vanishing gradients:**
 - ▶ Gradients become smaller and smaller and can become close to 0.
 - ▶ Can slow down or stop the learning process (very small weights' update).
- ▶ **Exploding gradients:**
 - ▶ Gradients become larger and larger as backpropagation progresses.
 - ▶ Learning can become unstable (large weights update) and diverge.

Design issues : Vanishing gradients and activation functions

As the number of layers goes up, the gradient is more likely to vanish during backpropagation.

- ▶ Using the ReLU activation function instead of tanh or sigmoid units can reduce this problem since its gradient does not go to zero as the input goes to zero.
- ▶ The Sigmoid and Tanh functions saturates at 0 or 1 when inputs become small or large.



Source: <https://towardsdatascience.com/why-rectified-linear-unit-relu-in-deep-learning-and-the-bestpractice-to-use-it-with-tensorflow-e9880933b7ef>



Existing activation functions

Summary of existing activation functions:

[https://ml-
cheatsheet.readthedocs.io/en/latest/activation_functions.html](https://ml-cheatsheet.readthedocs.io/en/latest/activation_functions.html)



Jupyter Notebook

Jupyter Notebook

Neural Network Design Issues
Coding Example



Advantages v.s. Disadvantages

Disadvantages

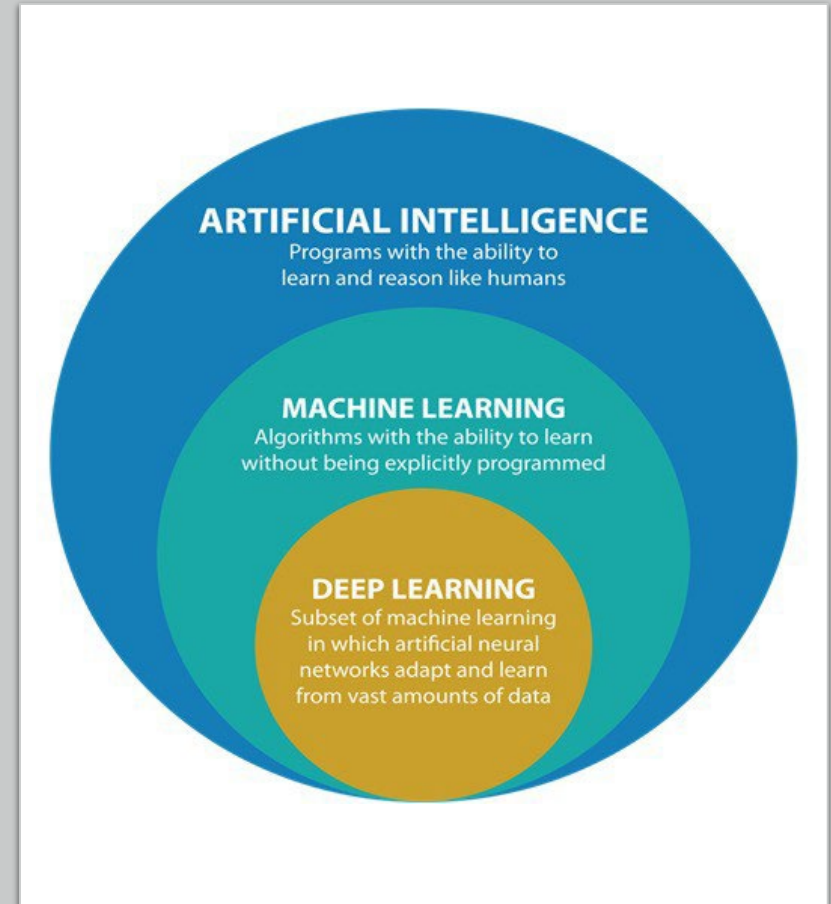
- Long training time
- Require to empirically determine, e.g., the network topology or “structure.”
- Difficult to interpret the symbolic meaning behind the learnable weights and hidden units in the network

Advantages

- High tolerance to noisy data
- Widely and empirically successful on real- world data, e.g., hand-written letters
- Algorithms are inherently parallel
- Techniques have recently been developed for the extraction of rules from trained neural networks
- Deep neural networks are powerful

SUMMARY

- Neural Nets
 - Multilayer Perceptron Architecture
 - Nonlinear Activation Functions
 - Training: Backpropagation algorithm
- Design Issues
 - Evaluation
 - Regularization techniques
 - Learning Rate
 - Initialization
 - Vanishing Gradient Problem
 - Tips ...





Resources

- Coding Libraries/Practice

- Python Machine Learning (3rd Edition) by Sebastian Raschka at <https://github.com/rasbt/python-machine-learning-book-3rd-edition>
- <https://playground.tensorflow.org/>

- Book Chapters

- Chapter 6.7, 6.8 Introduction to Data Mining by Kumar et al.
- <https://www.deeplearningbook.org/> Part II Deep Networks, chap. 8 and 11.

- Others

- <https://github.com/daiwk/ml-yearning/blob/master/Ng-MLY01-13.pdf>
- <http://karpathy.github.io/2019/04/25/recipe/>